

Lux Gives Terminal-Based AI Coding Agents a Visual Canvas

For the agent: describe UI as a typed JSON element tree, and Lux renders it natively with GPU-accelerated Dear ImGui on the developer's own machine, with no vendor compute

For the developer: see complex agent output as filterable tables, dashboards, and diagrams instead of scrollbar, and drive a terminal-based agent through the GUI, from a fixed, curated element set that composes without limit

Press Release

AUSTIN, TX, January 2027. Punt Labs today released Lux, a local-first visual output surface for terminal-hosted AI coding agents. The most capable agentic coding tools, Claude Code[1] and Codex[2], run in the terminal on the developer's own machine, but their structured output is trapped in scrollbar: a table, a dependency diagram, or a metrics dashboard arrives as monospace text or an ad hoc screenshot. Lux gives those agents a screen. An agent describes what it wants as a tree of typed JSON elements; Lux renders it with Dear ImGui[3] in a persistent, GPU-accelerated window, entirely on the developer's machine, with no vendor cloud. Unlike web dashboards that require running the agent on vendor compute, or the terminal that has no structure at all, Lux keeps the agent local and gives its output a surface the developer can read, filter, and click.

Problem

The most capable coding agents, Claude Code and Codex, live in the terminal: full tool use, MCP integration, multi-turn autonomous operation, and output that is all text. That power comes with no graphical surface. When an agent produces a table of failing tests, a diagram of a module graph, or a dashboard of build status, the developer reads it as raw scrollbar. There is no way to filter the table, select a row, or keep the view live as the underlying data changes. The most powerful agents run exactly where there are no graphics.

A second problem sits on top of the first: developers rarely run just one agent. Several run at once, across repositories, sessions, and machines, and tracking what each is doing is hands-on and error-prone. This is an aggregation problem: many independent agents, one developer, and no shared place to watch them.

Vendors are answering with hosted web dashboards that often require running the *entire agent* (not just the model call) on the vendor's compute: the developer trades local control for a viewport. IDE-hosted tools and MCP-host apps solve pieces of this but not the whole problem: one manages agents with text-only output, the other renders rich interfaces but cannot drive a terminal agent's autonomous, multi-turn workflow end-to-end (see [FAQ 2](#)). The result is a split: the tools with the most agentic power have no visual output, and the tools with visual output lack the agentic power.

Solution

Lux is a persistent local service that agents write to over MCP. An agent calls a tool such as `show_table` or `show_dashboard` with structured content; the display renders it in a native GPU-accelerated window using Dear ImGui. Multiple agents share the one window, each in its own frame with clear attribution, so a developer glances at a single surface and sees what every session is doing. Everything runs on the developer's machine: the Hub holds authoritative UI state, the display renders a replica, and no data leaves the box.

Lux is not a passive canvas. It ships two capabilities that matter in daily work. First, a rich data-display core: filterable, searchable tables with row selection and detail panels; dashboards that compose metric cards, charts, and tables; diagrams and images. The element vocabulary is deliberately limited: a curated, fixed set of built-in kinds that compose without limit, the same way a small set of ImGui widgets does. This is the core idea. Lux is not an IDE for agents; it prescribes no screens and no workflows. It is a medium through which an agent composes whatever UI it needs, so agents communicate through arbitrary interfaces built from one fixed vocabulary. Second, and the direction Lux is heading, an ask-the-user interaction loop: an agent can pose a question through a dialog, checkbox, or button and receive the answer back. When the user clicks, the interaction routes back to the owning Hub, the real handler runs there, and the Hub re-sends the updated view. That loop is the first interaction primitive, the seed of a two-way surface where the user drives a terminal-based agent through a GUI instead of only reading its output. The canonical example is the beads issue browser: a filterable table that auto-refreshes on every `bd` command through a post-tool hook, running live in every active agent session (see [FAQ 8](#)).

Customer Quote

"I keep four Claude Code sessions open across three repos. Before Lux I had four terminals and no idea what any of them were doing without scrolling back. Now one window shows me a beads board, a test-run table, and a build dashboard, one frame per agent, all live. When a migration script needed me to confirm a destructive step, the agent didn't dump a wall of text and wait; it put a dialog on the screen with the two options. I clicked. That is the difference between watching a transcript and using a tool."

— **Maya Torres**, Staff Engineer at a Series B DevTools Startup

Getting Started

Getting started takes one command:

1. Install:
`curl -fsSL https://raw.githubusercontent.com/punt-labs/lux/main/install.sh | sh`
2. Restart Claude Code. The Lux display window opens automatically the first time any agent sends visual output.
3. Ask an agent to show you something: "show me the beads board," "put the test results in a table." The view appears in the window, live.

No API key is required. Lux runs on macOS and Linux, needs Python 3.13+, and installs the `punt-lux` package plus the Claude Code plugin.

Spokesperson Quote

"The most powerful coding agents live in the terminal, and the terminal gives them one output mode: text. We built Lux so an agent can hand the developer a real surface (a table you filter, a dashboard you scan, a dialog you click) without moving the agent onto someone's cloud. The part I care most about is the interaction loop: the agent asks, the user answers,

and the answer flows back to the handler that asked. That turns a display into a conversation. And because the whole thing is local and native, 'Lux is showing me stale data' is not a failure mode: the Hub owns the state and re-sends the view when it changes."

— **Jim Freeman**, Punt Labs

Call to Action

Lux is free and open-source under the MIT license. Visit <https://github.com/punt-labs/lux> to install, read the target architecture, or contribute a new element kind.

Frequently Asked Questions

External FAQs

Q1: What is Lux and who is it for?

Lux is a visual output surface for terminal-hosted AI coding agents. It is a persistent, GPU-accelerated window, built on Dear ImGui[3], that agents write to and read from. It is for developers who run terminal agents such as Claude Code[1] or Codex[2] and want their structured output as something they can see and act on: filterable tables, dashboards, diagrams, and interactive prompts, rather than scrollbar text. The element kinds are fixed and curated; they compose without limit, covering any layout the way a small set of ImGui widgets does. Everything renders and runs on the developer's own machine. Lux runs on macOS and Linux.

Q2: How is this different from terminal scrollbar, Claude Desktop, or vendor dashboards?

Terminal agents (Claude Code, Codex) are the most capable but output only text. Claude Desktop renders rich interfaces and is itself an MCP host, but it cannot drive a coding session end-to-end: it lacks the terminal agent's autonomous, filesystem-level, multi-turn tool-driven workflow. Vendor web dashboards provide visuals but often require running the whole agent on vendor compute. IDEs such as Cursor manage agents with text-only output.

Lux fills the gap: a native visual surface for terminal-hosted agents, running on the developer's machine. It is not a chat interface, not a cloud console, and not an IDE for agents. Lux prescribes no screens or workflows; it lets agents communicate through arbitrary interfaces of their own composition. It is a screen the agents share: one window, one frame per agent, live. See [FAQ 9](#) for the detailed competitive analysis.

Q3: How do I get started?

One command installs Lux and registers the Claude Code plugin. Restart Claude Code. The display window opens by itself the first time an agent sends visual output. There is no separate app to launch and no configuration to edit. Ask an agent to show you a table or a board and it appears.

Q4: How much does it cost?

Free and open-source under the MIT license. There is no paid tier, no account, and no API key requirement for the display itself. Lux consumes the same agent you already run; it adds no per-call cost of its own.

Q5: What happens to my data?

Everything stays local. Agents talk to the Hub over MCP; the Hub talks to the display over a Unix domain socket under `/tmp/lux-<user>/`. Per-project configuration lives in your repository at `<repo>/punit-labs/lux.md`. There is no telemetry and no cloud dependency: nothing about your code, data, or agents' output leaves the machine. The Hub owns authoritative UI state and the display renders a replica; render calls never cross the wire, only UI state does.

Q6: What can Lux render today, and can agents ask me questions?

Two things ship and are in daily use. The data-display core covers filterable and searchable tables with row selection and detail panels, dashboards (metric cards, charts, tables), diagrams, and images: 25 element kinds in all. The interaction loop lets an agent ask you a question through the UI: a dialog, a checkbox, or a button, with your answer routed back to the agent's handler. That ask-the-user loop shipped with the Hub/Display interaction model. Sessions and apps also own entries in the shared menu bar (and the matching World menu): each entry belongs to the live session that registered it, and a click is serviced instantly by that session's own running process. Completing coverage of every ImGui widget as a built-in element is the v2 direction (see [FAQ 13](#)).

Internal FAQs

Value & Market

Q7: What is the total addressable market?

The primary market is developers running terminal-hosted AI coding agents: Claude Code[1] and Codex are the leading tools, both terminal-hosted. That population is a subset of professional developers, expanding as agentic coding moves from novelty to daily practice. We decline a precise headcount: no reliable public figure for active terminal-agent developers exists to cite without inventing it. As an order-of-magnitude *assumption* only (not a measurement), the cohort is plausibly 10^5 – 10^6 today, a small fraction of the world’s professional developers, and growing. Positioning, not sizing, is the point. Lux is free and MIT-licensed; its value to Punt Labs is ecosystem pull, not license revenue (see [FAQ 15](#)). The relevant question is not “how many will pay” but “how many terminal-agent developers would keep a local visual surface open once they have one.” That is the hypothesis current internal use supports and external use has not yet tested.

Q8: What evidence do we have that customers want this?

Lux is released (v0.22 on PyPI as `punt-lux` and on the plugin marketplace) and used daily. The precise scope matters and we state it plainly: the daily user is *one* primary human developer (Jim Freeman) plus the agent fleet he runs across the Punt Labs repositories. When we say the beads issue browser (a filterable table that auto-refreshes on every `bd` command through a post-tool hook) “runs in every session,” that counts *agent* sessions, not distinct human users. It is a genuine daily-driver signal, but $n=1$ human. That single view still demonstrates the two claims the product rests on: the data-display core has real daily value, and the live-refresh hook makes the surface feel current without the agent having to coordinate updates. The interaction loop, an agent asking a question and receiving the answer through a dialog, is exercised in real sessions, not just tests. A second proof point now exists at the application level: `vox` (the companion voice tool) ships a music player as a Lux Hub app – its daemon registers a menu entry, receives the user’s clicks over its persistent connection, and drives playback – accepted end-to-end in a live demo (a pixel click producing audio). Still Punt Labs’ own tool, so still internal evidence, but it validates the app surface, not just the agent one.

The honest gap: this evidence is *internal*. The target customer (an external developer running several terminal agents at once, the profile in the press-release quote) has not used Lux yet. Punt Labs is the current *user* and proof point, not the market. External adoption is zero: no usage data, no retention figures, no third-party testimonials. Validation confined to the company that built the product is, in the four-risks framing, a High Value-risk signal, and this document rates it as such. Every “value” claim here should be read against that boundary (see [FAQ 10](#) and the Value risk in the Risk Assessment).

Q9: Who are the competitors and why will we win?

The landscape divides on two axes: agentic power and visual surface.

- **Claude Code / Codex.** Terminal-hosted, the most capable agentic tools: full tool use, MCP, multi-turn autonomy. No visual output beyond scrollbar. *These are Lux’s integration targets, not rivals.*
- **Claude Desktop.** Renders novel visual interfaces and is the canonical MCP host, but lacks the terminal agent’s autonomous, filesystem-level, multi-turn workflow end-to-end.

- **Cursor.** IDE with strong agent management, but structured output stays text. Extensible via VS Code extensions, which are IDE-scoped and hand-written.
- **Vendor web dashboards.** Provide visual management, but often require running the whole agent on vendor compute. The developer gives up local control for a viewport.
- **Jupyter.** Rich visual output via the display protocol, but notebook-first, not agent-first.

Primary threat: Anthropic-native visual output. Anthropic has every incentive to build visual panels into Claude Code. A basic built-in display would narrow Lux's advantage to "native GPU-accelerated and fully local vs. vendor-native." Survival rests on three structural facts a competitor cannot copy cheaply: Lux renders natively with ImGui on the developer's GPU and keeps the agent local; Lux is a shared surface across *multiple* agents and tools, not a panel bolted onto one vendor's client; and the Hub/Display split lets an agent on a remote box drive a GUI on your local screen, because UI state rather than render calls crosses the wire (see [FAQ 12](#)). A web-native competitor would have to abandon web-only rendering to match the local-native path: a platform decision they are unlikely to reverse.

Q10: What is the next step to validate external demand?

The riskiest assumption is that Lux's internal value transfers to developers outside Punt Labs. The first external-validation experiment is to put Lux in front of external developers who run 2+ concurrent agent sessions daily and measure whether they keep the window open past the first day. Concretely: recruit a cohort, ship them the one-command install, and instrument nothing on their machine (no telemetry by design); instead, interview at day 7. The decision threshold: if fewer than half describe the display as something they would miss if it were removed, revisit the positioning before investing in the v2 element-coverage buildout.

Q11: How do external developers discover an MIT-licensed tool?

Three concrete channels, none of them yet proven for Lux.

- **The Claude Code plugin marketplace.** Lux is published there; a developer browsing plugins can find and install it. This is the most direct channel and requires no separate discovery of Punt Labs.
- **GitHub.** The repository is public and MIT-licensed; the one-command install lives in the README. This is organic search and word-of-mouth: real, but slow and unmeasured.
- **Cross-surfacing from the companion suite.** A developer who adopts vox, biff, or quarry encounters Lux as the surface those tools render into, and vice versa. The suite is the distribution flywheel the ecosystem thesis depends on (see [FAQ 15](#)).

The honest position: distribution is currently a hope, not a proven motion. "Organic and word-of-mouth" is not a strategy, and we have no acquisition data. The marketplace listing and the suite cross-surfacing are the two levers we can pull; whether either converts external installs is exactly what the day-7 cohort (see [FAQ 10](#)) begins to answer.

Technical

Q12: What is the architecture?

Lux is a three-tier distributed Hub/Display system. The target design lives in `docs/architecture/target/target.md`.

- **luxd**: the session Hub and the engine. It decodes wire elements, holds authoritative UI state, ownership, and handler dispatch in `HubDisplay`, pushes element replicas to the display, re-dispatches interactions that come back from clicks, and serves every client surface – MCP (streamable HTTP), REST, and the CLI – through one typed operations facade, so a capability runs the same engine code whichever surface it entered from.
- **lux-display**: the ImGui renderer. It stores a full copy of the UI, renders it every frame, and wraps handlers so that a click forwards back to the owning Hub instead of executing locally.
- **lux mcp-serve**: the per-session MCP server the Claude Code plugin spawns. It forwards the session's MCP traffic to luxd verbatim and holds a persistent listening connection on a background thread, which is what lets the session own menu entries whose clicks are answered instantly – by a running process, never by a turn of the model.

MCP is a first-class front door into the Hub, not a thin adapter. The boundary rule is strict: UI state crosses the wire (scene trees, serialized element objects, interaction messages); render calls do not. Interactive elements use two-tier handler dispatch (the Hub owns the real handler and the Display forwards the interaction back), so the Hub and Display can be separate processes, potentially on separate machines.

This lets the tiers sit on different machines. The Display connects to the Hub over a Unix socket; agents and apps connect to the Hub over HTTP and WebSocket on localhost (127.0.0.1:8430) by default. A remote agent works today by tunneling that localhost port over SSH. The practical payoff: your most powerful agent can run on a remote box and still drive a GUI on your local screen, because only UI state crosses the wire.

Q13: What is the v1-to-v2 roadmap, and what is the migration risk?

v1 (shipped): the agent shows the developer complex data and visualizations. The rich data-display core, plus the ask-the-user loop (dialog, checkbox, button) as the first interaction primitive, on the Hub/Display model.

v2 (direction): the developer interacts with the agent through a GUI. The surface turns two-way, so the user drives a terminal-based agent through a GUI instead of only viewing its output. Completing ImGui widget coverage is the means to that end, not the end in itself: richer interaction needs more input primitives (sliders, inputs, selects) and the composites that arrange them. This is still Punt Labs building one fixed, curated vocabulary, explicitly *not* agents authoring arbitrary extensions and not a plugin registry, because a bounded element set already composes without limit.

The element-migration risk this section used to carry is retired: all 25 element kinds are on the distributed Element-ABC path with two-tier remote handler dispatch, and the legacy render path is deleted. The remaining v2 feasibility exposure is the rendering layer's test gap – the renderer is verified by eyes and introspection tools, not by automated visual regression – and the interaction vocabulary itself: richer GUI-driven work needs more input primitives and the composites that arrange them, each a protocol-plus-renderer change with an evaluator gate (see the Risk Assessment).

Q14: If Lux succeeds, what breaks first?

Three known limits, stated with their current status rather than reassurance.

1. **The image texture cache grows without eviction.** Every image an agent shows uploads a texture that is never freed: an unbounded memory leak under sustained image use. Status: known and tracked only as a Should-Do (a texture-cache eviction policy, [Feature 13](#)); it is a real limitation today, not a solved problem. It is the first thing that breaks under a long-running session that renders many distinct images.
2. **Renderer maintainability debt.** `display/server.py` is roughly 1,500 lines, well over the 300-line target. Decomposition is in progress (the original monolith was split once; the legacy renderer modules died with the element migration). This does not break at runtime, but it slows every rendering change and is why the render layer is undertested. Status: actively being decomposed, incrementally.
3. **Many-agent contention on one window.** The shared window is a single surface; a large agent fleet all pushing scenes contends for one display and one Hub. The whole-tree resend replication policy is simple but re-sends the whole affected UI on change. Status: adequate at the current fleet size; a heavy multi-agent load is untested and would be the first place the simple replication model needs a diff protocol.

Business**Q15: What is the revenue model?**

Lux has no direct revenue. It is free and MIT-licensed. Its strategic value is ecosystem pull for the Punt Labs agent-tool suite. Lux is the visual counterpart to the other tools in the suite: vox (voice), biff (presence), and quarry (search). It surfaces them, too: the beads board, presence panels, and status views render in the Lux window. A developer who adopts Lux is one step from the rest of the suite. On the cost side there is no COGS (Lux runs entirely on the developer's own machine), so Punt Labs' development and maintenance time is the only real expense. This is an infrastructure and ecosystem play, acceptable because it is treated as a public good with strategic return, not a profit center. If Lux does not pull adoption of the companion tools, the investment is a public good without business return: that is the viability risk.

Q16: What are the key metrics for success?

1. **External retention:** of developers who install Lux outside Punt Labs, the share still running it after one week. This is the metric that converts the internal-only evidence into external validation.
2. **Companion-tool pull:** share of Lux installers who adopt at least one other Punt Labs tool (biff, vox, quarry). This measures the ecosystem thesis directly.
3. **Hub-app adoption:** applications built on the Hub's client surface beyond the agent tools (one today: vox's music player), tracking whether the app surface pulls its own builders.
4. **Interaction usage:** sessions that use the ask-the-user loop and the menu surface, distinguishing Lux-as-display from Lux-as-interaction-surface.

Decision threshold: if external one-week retention stays near zero after a genuine external trial, the product is an internal tool, not a market one, and the roadmap should stop treating external adoption as a near-term goal.

Q17: Why now? What has changed?

Three things converged. First, terminal-hosted agents, Claude Code and Codex, became the most capable agentic coding tools, and developers now run several at once, all producing output invisible beyond scrollbar. Second, the vendor response is web dashboards that pull the agent onto remote compute, opening a clear gap for a local-first visual surface. Third, Lux itself proved the display and the interaction loop have real daily value internally, and the Hub/Display architecture that makes the shared, attributed, multi-agent window possible is shipped, not speculative.

Q18: What are we explicitly not building?

Lux is a curated visual surface, not a platform for arbitrary code. It will not become a self-extending display where agents author and register their own extensions; there is no extension registry, no community extension store, and no in-display LLM authoring agent. Its client surfaces (MCP, REST, CLI, library) bind to the developer's own machine: there is no hosted service, no public web API, and no account. It has one rendering mode (native ImGui), not several. It does not host local model inference, does not render to the web or Electron, and does not target mobile or tablet. These exclusions are what keep the built-in vocabulary coherent and the trust model simple; the full list is in the Feature Appendix. One earlier exclusion was reversed: applets are now in scope as *Lux applets* – programs on the library client that own menu entries and service clicks (the vox music player is the first third-party one; lux-beads the first bundled one).

Risk Assessment

Risk	Rating	Assessment
Value	High	All demonstrated value is <i>internal</i> : the current users are one Punt Labs developer plus his agent fleet (see FAQ 8), not the target customer the product is built for: external developers running several terminal agents. “Runs in every session” counts agent sessions, not distinct humans; external adoption is zero. Validation confined to the company that built Lux is a High Value-risk signal in the four-risks framework, so this row is rated High, not Medium. It is why FAQ 10 makes the day-7 external trial the single most important next step.
Usability	Medium	Install is one command and the display opens automatically on first output, so the intended onboarding is simple. But that flow is proven <i>internally</i> only; it has never completed on a machine outside Punt Labs. Real prerequisites add friction: Python 3.13+, macOS/Linux only, and a curl-pipe-to-sh install that asks the user to trust a script. The day-7 external cohort (see FAQ 10) is the first real measurement of onboarding-completion; until then the Low-friction claim is a design intent, not a measured result.
Feasibility	Medium	The core is shipped and works: Hub/Display, the MCP and REST front doors, the data-display core, the interaction loop (dialogs, buttons, session menus with instant click servicing), and formally-verified lifecycle and atomicity models. The element migration that used to be the primary exposure is complete: 25 of 25 kinds on the Element-ABC path, legacy path deleted. The remaining exposure is the rendering layer’s automated-test gap (correctness there is verified by introspection tools and human demos, not visual regression in CI) and the ordinary risk of growing the interaction vocabulary (see FAQ 13).
Viability	Medium	MIT-licensed with no direct revenue. Strategic value depends entirely on ecosystem pull for the companion Punt Labs tools. If Lux does not drive adoption of biff, vox, and quarry, it is a public good without business return. Acceptable only if treated as infrastructure, which is the stated intent (see FAQ 15).

Feature Appendix

Defines the scope boundary for Lux. Must Do is the shipped v1 identity; Should Do is the v2 direction; Won't Do is the firm exclusion set that keeps the built-in vocabulary coherent and the trust model simple.

Must Do

The core that makes Lux what it is. These ship in v1.

- F1. Typed JSON element protocol:** agents describe UI as a tree of typed elements: tables, text, plots, groups, buttons, and more; the protocol is the API surface
- F2. Native GPU-accelerated rendering:** a persistent, shared ImGui window painting every frame on the developer's own GPU, with no web stack and no vendor compute
- F3. Hub/Display io-model:** luxd holds authoritative UI state; the display renders a replica; UI state crosses the wire, render calls do not
- F4. Data-display core:** filterable, searchable tables with row selection and detail panels; dashboards; diagrams; images; 25 element kinds in all
- F5. Ask-the-user interaction loop:** an agent poses a question through a dialog, checkbox, or button and receives the answer back via Hub-side handler dispatch
- F6. Two-tier remote handler dispatch:** the Hub owns handlers; the Display forwards interactions back and the real handler runs on the Hub's copy
- F7. Multi-agent shared window:** multiple agents share one window, each in its own frame with clear attribution
- F8. MCP front door:** a first-class entry point: luxd serves MCP over streamable HTTP, and each Claude Code session runs `lux mcp-serve`, whose persistent connection lets the session own menu entries with instantly-serviced clicks
- F9. Introspection and control surface:** `inspect_scene`, `list_scenes`, `screenshot`, and related tools let an agent verify what actually rendered
- F10. One-command install with auto-open display:** a single curl command installs the package and plugin; the window opens on first output
- F11. Formally-verified singleton lifecycle:** the spawn/reap/bind concurrency (the display singleton lifecycle only, not rendering, protocol, or dispatch) is specified in Z and ProB-model-checked for singleton-serving, never-unlink-live, no-two-winners, clean-exit, and deadlock-freedom

Should Do

The v2 direction. Valuable, not v1-blocking.

- F12. Complete ImGui element coverage:** every ImGui widget available as a curated, built-in element; the point is not coverage for its own sake but the input primitives that enable rich GUI-driven interaction with the agent
- F13. Texture-cache eviction policy:** bound the image texture cache, which today grows without eviction
- F14. Companion-tool surfacing:** render biff presence panels, vox controls, and quarry search as first-class Lux views alongside the beads board

Won't Do

Explicitly excluded to keep the vocabulary curated and the trust model simple.

- F15. Self-extending extension registry:** agents will not author and register arbitrary display extensions at runtime; the fixed, curated element set already composes without limit, so infinite element kinds are unnecessary, and Lux stays a built-in vocabulary rather than a plugin platform
- F16. Community extension store:** no marketplace of third-party extensions with ratings and verification badges, which sit outside the trust model Lux commits to
- F17. In-display LLM authoring agent:** Lux does not embed a model that writes new UI code inside the display
- F18. Hosted service API:** the MCP, REST, CLI, and library surfaces are thin local clients of one engine on the developer's machine; Lux does not become a hosted web service, expose a public API, or require an account
- F19. Multiple rendering modes:** one native ImGui mode only, with no texture-buffer or separate-OS-window rendering modes
- F20. Applets as arbitrary in-display applications:** what IS in scope (reversing this row's earlier full exclusion, by operator ruling 2026-08-01): *Lux applets* – external programs on the library client that own menu entries and render through the Hub, either shipped by their own tool (vox's music player) or bundled with Lux for software that is not Punt Labs-built (*lux-beads*; *top/clock/irc-class* candidates). What stays out: applications running *inside* the display process, and Lux itself becoming an application platform with in-display code execution
- F21. Local model inference:** Lux renders; it does not host models
- F22. Web or Electron rendering:** Lux renders natively via ImGui on the developer's GPU. HTML, Electron, and webviews are out of scope, since web rendering leads to the vendor-compute dependency Lux exists to avoid
- F23. Mobile or tablet:** desktop only (macOS, Linux)

References

- [1] Anthropic. *Claude Code*. Agentic coding tool for the terminal. 2025. URL: <https://docs.anthropic.com/en/docs/agents-and-tools/claude-code/overview>.
- [2] OpenAI. *Codex CLI*. Terminal-hosted agentic coding tool. 2025. URL: <https://github.com/openai/codex>.
- [3] Omar Cornut. *Dear ImGui*. Immediate-mode GUI library for C++. 2025. URL: <https://github.com/ocornut/imgui>.